

Settings

Ollama model

llama3.2

☐ Use Ollama

Top-k evidence chunks

5

Enterprise AI Knowledge Agent Platform

Local/private RAG + SQL agent + image notes + confidence scoring + hallucination checks

Project Question

Can a private enterprise AI agent platform answer questions across documents, SQL tables, and image notes with citations, confidence scoring, and hallucination checks?

[Ask Agent](#) [Evaluation](#) [Architecture](#)

Ask a question

Which portfolio projects have the highest priority and why?

Run Agent

Agent Route

sql_agent

Answer

Ollama disabled. Retrieved evidence summary:\n\nproject_portfolio_metrics.csv chunk sql: project domain primary_skill risk_level portfolio_priority CAPSDAC Forecasting Public Sector Education Forecasting and data governance High 5 Local RAG Assistant Enterprise AI RAG and local LLM deployment Medium 5 Medical Imaging

Citation Confidence

```
▼ {  
  "confidence_score" : 0.75  
  "confidence_label" : "high"  
  "reason" : "Average retrieval score=1.000; distinct sources=1."  
}
```

Hallucination Risk

```
▼ {  
  "risk_level" : "low"  
  ▶ "flags" : []  
}
```

Evidence

	filename	chunk_index	score	modality	preview
0	project_portfolio_metrics.csv	sql	1	sql	project domain pri

SQL Query

```
SELECT project, domain, primary_skill, risk_level, portfolio_priority  
FROM portfolio_metrics  
ORDER BY portfolio_priority DESC, project ASC  
LIMIT 10
```

	project	domain	primary_skill	risk_level	port
0	CAPSDAC Forecasting	Public Sector Education	Forecasting and data governance	High	
1	Local RAG Assistant	Enterprise AI	RAG and local LLM deployment	Medium	
2	Medical Imaging Clinical AI	Healthcare	Multimodal ML and explainability	High	
3	GenAI Drug Discovery	Healthcare / Life Sciences	Generative modeling and evaluation	Medium	
4	Transportation Telemetry	Transportation	Telemetry lakehouse and neural prediction	Medium	