

Settings

Ollama model

llama3.2

☐ Use Ollama

Top-k evidence chunks



Enterprise AI Knowledge Agent Platform

Local/private RAG + SQL agent + image notes + confidence scoring + hallucination checks

Project Question

Can a private enterprise AI agent platform answer questions across documents, SQL tables, and image notes with citations, confidence scoring, and hallucination checks?

Ask Agent **Evaluation** Architecture

Run evaluation first:

```
python evaluation/evaluate_agent.py
```

Evaluation Summary

	hit_rate	average_confidence	n_questions
0	1	0.565	5

Evaluation Results

	question	expected_source_keyword	route	
0	What should healthcare AI governance include?	healthcare_ai_governance	document_rag_agent	
1	Why is de-identification important for public-sector analytics?	public_sector_analytics	document_rag_agent	
2	What is an enterprise RAG architecture?	enterprise_rag_architecture	document_rag_agent	
3	Which portfolio projects have highest priority?	project_portfolio_metrics	sql_agent	
4	What does the platform architecture diagram show?	platform_architecture_note	image_agent	